

Navya_Bhat_HW6

November 1, 2023

```
[1]: from pathlib import Path
import sys

if 'google.colab' in str(get_ipython()):
    from google.colab import drive
    drive.mount('/content/drive')

    base_folder = Path('/content/drive/MyDrive/')
    data_folder = Path('/content')

    !pip install pytorch-lightning==2.0.9 -qq
    !pip install torchmetrics -U -qq
    !pip install fastdownload -U -qq
    !pip install fastai -U -qq
    !pip install wandb -U -qq

else:
    base_folder = Path('/home/harpreet/Insync/google_drive_shaannoor')
    data_folder = Path('/home/harpreet/data')
```

Mounted at /content/drive

727.7/727.7

kB 9.9 MB/s eta 0:00:00

805.2/805.2

kB 41.6 MB/s eta 0:00:00

2.1/2.1 MB

21.6 MB/s eta 0:00:00

190.6/190.6

kB 22.3 MB/s eta 0:00:00

243.2/243.2

kB 27.5 MB/s eta 0:00:00

Preparing metadata (setup.py) ... done

62.7/62.7 kB

8.3 MB/s eta 0:00:00

Building wheel for pathtools (setup.py) ... done

```
[2]: custom_function_folder = base_folder/'data/custom-functions/fall_2023'
sys.path.append('/content/drive/MyDrive/')
model_folder = base_folder/'data/models/dl_fall_2023/fmnist/oct-31'
model_folder.mkdir(parents=True, exist_ok=True)
project_folder = base_folder/'data/fmnist'
project_folder.mkdir(parents=True, exist_ok=True)
```

```
[3]: # import Libraries
import yaml

import torch
import torchmetrics
from torchvision import transforms
import pytorch_lightning as pl
from pytorch_lightning import seed_everything
from pytorch_lightning.tuner import Tuner
from pytorch_lightning.callbacks import ModelCheckpoint, EarlyStopping, LearningRateMonitor
from pytorch_lightning.loggers import CSVLogger, WandbLogger
import wandb
import gc
import torch.nn as nn

from data_module_fmnist import FashionMNISTDataModule
from multiclass_lightning_module_v0 import MultiClassLightningModule
from shared_utils import plot_losses_acc
```

```
[4]: !pip show pytorch-lightning
```

```
Name: pytorch-lightning
Version: 2.0.9
Summary: PyTorch Lightning is the lightweight PyTorch wrapper for ML
researchers. Scale your models. Write less boilerplate.
Home-page: https://github.com/Lightning-AI/lightning
Author: Lightning AI et al.
Author-email: pytorch@lightning.ai
License: Apache-2.0
Location: /usr/local/lib/python3.10/dist-packages
Requires: fsspec, lightning-utilities, numpy, packaging, PyYAML, torch,
torchmetrics, tqdm, typing-extensions
Required-by:
```

```
[5]: class ResidualBlock(nn.Module):
    def __init__(self, in_channels, out_channels, stride=1):
        super(ResidualBlock, self).__init__()

        self.main_path = nn.Sequential(
```

```

        nn.Conv2d(in_channels, out_channels, kernel_size=3, stride=stride,
↪padding=1),
        nn.BatchNorm2d(out_channels),
        nn.ReLU(inplace=True),
        nn.Conv2d(out_channels, out_channels, kernel_size=3, padding=1),
        nn.BatchNorm2d(out_channels)
    )

    self.downsample = nn.Sequential()
    if stride != 1 or in_channels != out_channels:
        self.downsample = nn.Sequential(
            nn.Conv2d(in_channels, out_channels, kernel_size=1,
↪stride=stride),
            nn.BatchNorm2d(out_channels)
        )

    def forward(self, x):
        residual = x
        out = self.main_path(x)
        residual = self.downsample(residual)
        out += residual
        out = nn.ReLU(inplace=True)(out)
        return out

class SimpleResNet(nn.Module):
    def __init__(self, num_classes=10):
        super(SimpleResNet, self).__init__()

        self.model = nn.Sequential(
            nn.Conv2d(1, 16, kernel_size=7, stride=2, padding=3), # Output:
↪16x250x188
            nn.BatchNorm2d(16),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(kernel_size=3, stride=2, padding=1), # Output:
↪16x125x94
            ResidualBlock(16, 32, stride=2), # Output:
↪32x63x47
            ResidualBlock(32, 64, stride=2), # Output:
↪64x32x24
            ResidualBlock(64, 256, stride=2), # Output:
↪128x16x12
            nn.AdaptiveAvgPool2d((1, 1)) # Output:
↪256x1x1
        )

        self.fc = nn.Linear(256, num_classes)

```

```

def forward(self, x):
    x = self.model(x)
    x = x.view(x.size(0), -1)
    x = self.fc(x)
    return x

```

```

[6]: def count_parameters(model):
    total_params = sum(p.numel() for p in model.parameters())
    trainable_params = sum(p.numel() for p in model.parameters() if p.
        ↪requires_grad)
    return total_params, trainable_params

model = SimpleResNet(num_classes=10)
total_params, trainable_params = count_parameters(model)
print(f"Total parameters: {total_params}")
print(f"Trainable parameters: {trainable_params}")

```

Total parameters: 831914
Trainable parameters: 831914

```

[7]: #Function to load the model
def load_model(model_config):
    model = SimpleResNet(num_classes=10)
    return model

```

```

[8]: #function for transformations
def get_train_transforms(resize_height, resize_width, normalize_mean, ↵
    ↪normalize_std):

    return transforms.Compose(
        [
            transforms.Resize((resize_height, resize_width)),
            transforms.ToTensor(),
            transforms.Normalize(normalize_mean, normalize_std),
        ]
    )

def get_test_transforms(resize_height, resize_width, normalize_mean, ↵
    ↪normalize_std):

    return transforms.Compose(
        [
            transforms.Resize((resize_height, resize_width)),
            transforms.ToTensor(),
            transforms.Normalize(normalize_mean, normalize_std),
        ]
    )

```

```
[9]: trans1 = transforms.ToTensor()
# Transform 2: Normalize the tensor images.
# The specified mean and standard deviation values are dataset-specific.
trans2 = transforms.Normalize((0.2857,), (0.3528))

# Combine the above transformations into a single composite transform.
trans = transforms.Compose([trans1, trans2])
```

```
[10]: def load_datamodule(config, data_folder):
    # Fetch the correct transform function based on config and pass the
    ↪ appropriate arguments
    train_transform = get_train_transforms(**config['train_transform'])
    test_transform = get_test_transforms(**config['test_transform'])
    dm = FashionMNISTDataModule(
        data_dir=data_folder,
        train_transform=train_transform,
        test_transform=test_transform,
        **config['data_module']
    )
    return dm
```

```
[11]: # Load Lightning Module
def load_lightning_module(config, model):
    optimizer_cls = eval(config['optimizer_cls'])
    loss_fn = eval(config['loss_fn'])() # directly instantiate the loss
    ↪ function
    metric_cls = eval(config['metric_cls'])

    # If scheduler is defined, convert its string to class as well
    if config.get('scheduler_cls'):
        scheduler_cls = eval(config['scheduler_cls'])
        scheduler_options = config['scheduler_options']
        scheduler_params = config['scheduler_params']
    else:
        scheduler_cls = None

    lightning_module = MultiClassLightningModule(model=model,
                                                    optimizer_cls=optimizer_cls,
                                                    loss_fn=loss_fn,
                                                    metric_cls=metric_cls,
                                                    scheduler_cls=scheduler_cls,
                                                    ↪
    ↪ scheduler_options=scheduler_options,
                                                    ↪
    ↪ scheduler_params=scheduler_params,
                                                    **config['others']
    )
```

```
return lightning_module
```

```
[12]: # Load the trainer
def load_trainer(model, trainer_config, cl_config, batch_size, model_folder,
    logging=False, checkpointing=True, early_stopping=False):

    lr_monitor = LearningRateMonitor(**cl_config['lr_monitor'])
    callbacks = [lr_monitor]
    if checkpointing:
        model_checkpoint_callback = ModelCheckpoint(dirpath=model_folder/
    cl_config['log_dir'],
                                                    **cl_config['model_checkpoint'])
        callbacks.append(model_checkpoint_callback)

    if early_stopping:
        early_stop_callback = EarlyStopping(**cl_config['early_stopping'] )
        callbacks.append(early_stop_callback)

    if logging:
        # For WandB logger:
        wandb_logger = WandbLogger(project=cl_config['wandb']['project'],
    name=cl_config['wandb']['name'], save_dir=model_folder/cl_config['log_dir'])
        wandb_logger.experiment.config.update({'batch_size': batch_size,
    'epochs': trainer_config['max_epochs']})
        wandb_logger.watch(model)

        # For CSV logger:
        csv_logger = CSVLogger(save_dir=model_folder/cl_config['log_dir'],
    name=cl_config['csv']['name'])
        csv_logger.log_hyperparams(params={'batch_size': batch_size, 'epochs':
    trainer_config['max_epochs']})

        trainer = pl.Trainer(callbacks=callbacks,
                                logger=[csv_logger, wandb_logger],
                                **trainer_config)
    else:
        trainer = pl.Trainer(callbacks=callbacks,
                                **trainer_config)

    return trainer
```

```
[13]: # Function to load components
def load_components(model_config, data_module_config, lightning_module_config,
    data_folder, trainer_config,
    cl_config, batch_size, logging=False, checkpointing=True, early_stopping=False):
```

```

# Load the model
model = load_model(model_config)

# Load the data module
dm = load_datamodule(data_module_config, data_folder)

# Load the lightning module
lightning_module = load_lightning_module(lightning_module_config, model)

# Load the trainer
trainer = load_trainer(model, trainer_config, cl_config, batch_size,
↪model_folder, logging=logging,
                                checkpointing=checkpointing,
↪early_stopping=early_stopping)

return model, dm, lightning_module, trainer

```

```

[14]: def load_yaml(filepath):
        with open(filepath, 'r') as file:
            return yaml.safe_load(file)

```

```

[15]: # Load configurations from YAML files
def load_all_configs():
    model_config = load_yaml(project_folder/'model_config.yaml')
    data_module_config = load_yaml(project_folder/'data_module_config.yaml')
    lightning_module_config = load_yaml(project_folder/'lightning_module_config.
↪yaml')
    cl_config = load_yaml(project_folder/'callbacks_loggers_config.yaml')
    trainer_config = load_yaml(project_folder/'trainer_config.yaml')

    return model_config, data_module_config, lightning_module_config,
↪cl_config, trainer_config

```

```

[16]: def free_memory():
        """
        Attempts to free up memory by deleting variables and running Python's
↪garbage collector.
        """
        gc.collect()
        for device_id in range(torch.cuda.device_count()):
            torch.cuda.set_device(device_id)
            torch.cuda.empty_cache()
        gc.collect()

```

```

[17]: print(project_folder)

```

```

/content/drive/MyDrive/data/fmnist

```

```
[18]: # Running a training and validation batch
# Load components
free_memory()
seed_everything(42)
model_config, data_module_config, lightning_module_config, cl_config,
    ↪ trainer_config = load_all_configs()
# override default values
trainer_config['fast_dev_run']=True
model, dm, lightning_module, trainer = load_components(model_config,
    ↪ data_module_config,
                                                    lightning_module_config,
    ↪ data_folder, trainer_config,
                                                    cl_config,
    ↪ batch_size=data_module_config['data_module']['batch_size'],
                                                    logging=False,
    ↪ checkpointing=False, early_stopping=False)
dm.prepare_data()
trainer.fit(lightning_module, dm)
```

```
INFO:lightning_fabric.utilities.seed:Global seed set to 42
INFO:pytorch_lightning.utilities.rank_zero:GPU available: True (cuda), used:
True
INFO:pytorch_lightning.utilities.rank_zero:TPU available: False, using: 0 TPU
cores
INFO:pytorch_lightning.utilities.rank_zero:IPU available: False, using: 0 IPUs
INFO:pytorch_lightning.utilities.rank_zero:HPU available: False, using: 0 HPUs
INFO:pytorch_lightning.utilities.rank_zero:Running in `fast_dev_run` mode: will
run the requested loop using 1 batch(es). Logging and checkpointing is
suppressed.
```

```
Downloading http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/train-
images-idx3-ubyte.gz
```

```
Downloading http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/train-
images-idx3-ubyte.gz to /content/FashionMNIST/raw/train-images-idx3-ubyte.gz
```

```
100%|          | 26421880/26421880 [00:03<00:00, 7251610.57it/s]
```

```
Extracting /content/FashionMNIST/raw/train-images-idx3-ubyte.gz to
/content/FashionMNIST/raw
```

```
Downloading http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/train-
labels-idx1-ubyte.gz
```

```
Downloading http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/train-
labels-idx1-ubyte.gz to /content/FashionMNIST/raw/train-labels-idx1-ubyte.gz
```

```
100%|          | 29515/29515 [00:00<00:00, 137749.30it/s]
```

```
Extracting /content/FashionMNIST/raw/train-labels-idx1-ubyte.gz to
/content/FashionMNIST/raw
```



```

Downloading http://fashion-mnist.s3-website.eu-
central-1.amazonaws.com/t10k-images-idx3-ubyte.gz
Downloading http://fashion-mnist.s3-website.eu-
central-1.amazonaws.com/t10k-images-idx3-ubyte.gz to
/content/FashionMNIST/raw/t10k-images-idx3-ubyte.gz
100%|      | 4422102/4422102 [00:01<00:00, 2572010.57it/s]
Extracting /content/FashionMNIST/raw/t10k-images-idx3-ubyte.gz to
/content/FashionMNIST/raw

Downloading http://fashion-mnist.s3-website.eu-
central-1.amazonaws.com/t10k-labels-idx1-ubyte.gz
Downloading http://fashion-mnist.s3-website.eu-
central-1.amazonaws.com/t10k-labels-idx1-ubyte.gz to
/content/FashionMNIST/raw/t10k-labels-idx1-ubyte.gz
100%|      | 5148/5148 [00:00<00:00, 7961754.05it/s]
Extracting /content/FashionMNIST/raw/t10k-labels-idx1-ubyte.gz to
/content/FashionMNIST/raw

INFO:pytorch_lightning.accelerators.cuda:LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES:
[0]
INFO:pytorch_lightning.callbacks.model_summary:
  | Name          | Type                | Params
-----
0 | model          | SimpleResNet        | 831 K
1 | loss_fn        | CrossEntropyLoss    | 0
2 | train_metric   | MulticlassAccuracy  | 0
3 | val_metric     | MulticlassAccuracy  | 0
4 | test_metric    | MulticlassAccuracy  | 0
-----
831 K      Trainable params
0          Non-trainable params
831 K      Total params
3.328      Total estimated model params size (MB)
/usr/local/lib/python3.10/dist-
packages/pytorch_lightning/trainer/connectors/data_connector.py:442:
PossibleUserWarning: The dataloader, train_dataloader, does not have many
workers which may be a bottleneck. Consider increasing the value of the
`num_workers` argument` (try 8 which is the number of cpus on this machine) in
the `DataLoader` init to improve performance.
  rank_zero_warn(
/usr/local/lib/python3.10/dist-
packages/pytorch_lightning/trainer/connectors/data_connector.py:442:
PossibleUserWarning: The dataloader, val_dataloader, does not have many workers
which may be a bottleneck. Consider increasing the value of the `num_workers`
argument` (try 8 which is the number of cpus on this machine) in the

```

```

`DataLoader` init to improve performance.
rank_zero_warn(
Training: 0it [00:00, ?it/s]
Validation: 0it [00:00, ?it/s]
Epoch 1: Val_Loss: 2.29, Val_Metric: 0.14 |
INFO:pytorch_lightning.utilities.rank_zero:`Trainer.fit` stopped: `max_steps=1`
reached.
Train_Loss: 2.45, Train_Metric: 0.14

```

```

[19]: print(model_config, data_module_config, lightning_module_config, cl_config,
        ↪trainer_config)

```

```

{'num_features': 562500, 'hidden_dim1': 500, 'hidden_dim2': 500, 'num_classes':
10} {'train_transform': {'resize_height': 500, 'resize_width': 375,
'normalize_mean': [0.5], 'normalize_std': [0.5]}, 'test_transform':
{'resize_height': 500, 'resize_width': 375, 'normalize_mean': [0.5],
'normalize_std': [0.5]}, 'data_module': {'batch_size': 64, 'seed': 42}}
{'optimizer_cls': 'torch.optim.AdamW', 'loss_fn': 'torch.nn.CrossEntropyLoss',
'metric_cls': 'torchmetrics.Accuracy', 'scheduler_cls': 'None',
'scheduler_options': 'None', 'scheduler_params': 'None', 'others':
{'optimizer_params': {'weight_decay': 0}, 'num_classes': 10, 'learning_rate':
0.0001, 'log_every_n_steps': 1, 'log_test_metrics': True, 'display_metrics':
True}} {'log_dir': 'logs', 'lr_monitor': {'logging_interval': 'step'},
'model_checkpoint': {'monitor': 'val_metric', 'mode': 'max', 'save_top_k': 1,
'save_last': True}, 'early_stopping': {'monitor': 'val_metric', 'patience': 5,
'mode': 'max', 'verbose': True}, 'wandb': {'project': 'FMNIST', 'name':
'resnet'}, 'csv': {'name': 'csvlogger'}} {'max_epochs': 2, 'accelerator':
'auto', 'devices': 'auto', 'deterministic': False, 'log_every_n_steps': 1,
'gradient_clip_algorithm': 'norm', 'gradient_clip_val': 0, 'fast_dev_run': True,
'overfit_batches': 0.0, 'accumulate_grad_batches': 1, 'limit_train_batches':
1.0, 'limit_val_batches': 1.0, 'limit_test_batches': 1.0}

```

```

[20]: # Load components
free_memory()
seed_everything(42)
model_config, data_module_config, lightning_module_config, cl_config,
        ↪trainer_config = load_all_configs()
# override default values
trainer_config['max_epochs']=10
model, dm, lightning_module, trainer = load_components(model_config,
        ↪data_module_config,
                                                    lightning_module_config,
        ↪data_folder, trainer_config,
                                                    cl_config,
        ↪batch_size=data_module_config['data_module']['batch_size'],

```

```

                                                                    logging=False,
        ↪checkpointing=False, early_stopping=False)
dm.setup(stage='fit')
tuner = Tuner(trainer)
lr_finder = tuner.lr_find(lightning_module, datamodule=dm, min_lr=1e-5,
        ↪max_lr=1, num_training=30, mode='exponential')
fig = lr_finder.plot(suggest=True)
new_lr = lr_finder.suggestion()
print(new_lr)

```

```


INFO:lightning_fabric.utilities.seed:Global seed set to 42
INFO:pytorch_lightning.utilities.rank_zero:GPU available: True (cuda), used:
True
INFO:pytorch_lightning.utilities.rank_zero:TPU available: False, using: 0 TPU
cores
INFO:pytorch_lightning.utilities.rank_zero:IPU available: False, using: 0 IPUs
INFO:pytorch_lightning.utilities.rank_zero:HPU available: False, using: 0 HPUs
INFO:pytorch_lightning.utilities.rank_zero:`Trainer(limit_train_batches=1.0)`
was configured so 100% of the batches per epoch will be used..
INFO:pytorch_lightning.utilities.rank_zero:`Trainer(limit_val_batches=1.0)` was
configured so 100% of the batches will be used..
INFO:pytorch_lightning.utilities.rank_zero:`Trainer(limit_test_batches=1.0)` was
configured so 100% of the batches will be used..
WARNING:pytorch_lightning.loggers.tensorboard:Missing logger folder:
/content/lightning_logs
INFO:pytorch_lightning.accelerators.cuda:LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES:
[0]

Epoch 1: Val_Loss: 2.30, Val_Metric: 0.03 |
Finding best initial lr:  0%|          | 0/30 [00:00<?, ?it/s]

INFO:pytorch_lightning.utilities.rank_zero:`Trainer.fit` stopped: `max_steps=30`
reached.
INFO:pytorch_lightning.tuner.lr_finder:Learning rate set to 0.002154434690031884
INFO:pytorch_lightning.utilities.rank_zero:Restoring states from the checkpoint
path at /content/.lr_find_4c9b37df-b6d0-402f-8f32-bc0e0e36c765.ckpt
INFO:pytorch_lightning.utilities.rank_zero:Restored all states from the
checkpoint at /content/.lr_find_4c9b37df-b6d0-402f-8f32-bc0e0e36c765.ckpt

Train_Loss: 2.48, Train_Metric: 0.28
0.002154434690031884

```



```
[21]: free_memory()
seed_everything(42)
model_config, data_module_config, lightning_module_config, cl_config,
↳ trainer_config = load_all_configs()

# override default values
data_module_config['data_module']['batch_size']=128
lightning_module_config['others']['learning_rate']=0.002
trainer_config['max_epochs']=10
trainer_config['gradient_clip_val']=2
trainer_config['log_every_n_steps']=20

lightning_module_config['others']['optimizer_params']['weight_decay']=1
lightning_module_config['others']['learning_rate']=0.002
lightning_module_config['scheduler_cls']='torch.optim.lr_scheduler.
↳ ReduceLROnPlateau'
lightning_module_config['scheduler_params'] = {'mode': 'max', 'patience': 0,
↳ 'factor': 0.5, 'verbose': True}
lightning_module_config['scheduler_options'] = {'monitor': 'val_metric',
↳ 'interval': 'epoch', 'frequency': 1}
```

```

cl_config['lr_monitor']['logging_interval']='epoch'
cl_config['wandb']['project']='fminst'
cl_config['wandb']['name']='resnet'

model, dm, lightning_module, trainer = load_components(model_config,
↳data_module_config,
lightning_module_config,
↳data_folder, trainer_config,
cl_config,
↳batch_size=data_module_config['data_module']['batch_size'],
logging=True,
↳checkpointing=True, early_stopping=True)
dm.setup(stage='fit')
trainer.fit(lightning_module, dm)

```

INFO:lightning_fabric.utilities.seed:Global seed set to 42

<IPython.core.display.Javascript object>

wandb: Logging into wandb.ai. (Learn how to deploy a W&B server locally: <https://wandb.me/wandb-server>)

wandb: You can find your API key in your browser here:

<https://wandb.ai/authorize>

wandb: Paste an API key from your profile and hit enter, or press ctrl+c to quit:

.....

wandb: Appending key for api.wandb.ai to your netrc file:
/root/.netrc

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

wandb: logging graph, to disable use `wandb.watch(log_graph=False)`

INFO:pytorch_lightning.utilities.rank_zero:GPU available: True (cuda), used: True

INFO:pytorch_lightning.utilities.rank_zero:TPU available: False, using: 0 TPU cores

INFO:pytorch_lightning.utilities.rank_zero:IPU available: False, using: 0 IPUs

INFO:pytorch_lightning.utilities.rank_zero:HPU available: False, using: 0 HPUs

INFO:pytorch_lightning.utilities.rank_zero:`Trainer(limit_train_batches=1.0)` was configured so 100% of the batches per epoch will be used..

INFO:pytorch_lightning.utilities.rank_zero:`Trainer(limit_val_batches=1.0)` was

```

configured so 100% of the batches will be used..
INFO:pytorch_lightning.utilities.rank_zero:`Trainer(limit_test_batches=1.0)` was
configured so 100% of the batches will be used..
/usr/local/lib/python3.10/dist-
packages/pytorch_lightning/callbacks/model_checkpoint.py:617: UserWarning:
Checkpoint directory
/content/drive/MyDrive/data/models/dl_fall_2023/fmnist/oct-31/logs exists and is
not empty.
    rank_zero_warn(f"Checkpoint directory {dirpath} exists and is not empty.")
INFO:pytorch_lightning.accelerators.cuda:LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES:
[0]
INFO:pytorch_lightning.callbacks.model_summary:
  | Name          | Type                | Params
-----
0 | model          | SimpleResNet        | 831 K
1 | loss_fn        | CrossEntropyLoss    | 0
2 | train_metric   | MulticlassAccuracy  | 0
3 | val_metric     | MulticlassAccuracy  | 0
4 | test_metric    | MulticlassAccuracy  | 0
-----
831 K      Trainable params
0          Non-trainable params
831 K      Total params
3.328      Total estimated model params size (MB)
Sanity Checking: 0it [00:00, ?it/s]

/usr/local/lib/python3.10/dist-
packages/pytorch_lightning/trainer/connectors/data_connector.py:442:
PossibleUserWarning: The dataloader, val_dataloader, does not have many workers
which may be a bottleneck. Consider increasing the value of the `num_workers`
argument` (try 8 which is the number of cpus on this machine) in the
`DataLoader` init to improve performance.
    rank_zero_warn(

Epoch 1: Val_Loss: 2.30, Val_Metric: 0.04 |

/usr/local/lib/python3.10/dist-
packages/pytorch_lightning/trainer/connectors/data_connector.py:442:
PossibleUserWarning: The dataloader, train_dataloader, does not have many
workers which may be a bottleneck. Consider increasing the value of the
`num_workers` argument` (try 8 which is the number of cpus on this machine) in
the `DataLoader` init to improve performance.
    rank_zero_warn(

Training: 0it [00:00, ?it/s]

Validation: 0it [00:00, ?it/s]

Epoch 1: Val_Loss: 1.27, Val_Metric: 0.56 |

```

```


INFO:pytorch_lightning.callbacks.early_stopping:Metric val_metric improved. New
best score: 0.564
Train_Loss: 0.69, Train_Metric: 0.76
Validation: 0it [00:00, ?it/s]
Epoch 2: Val_Loss: 0.61, Val_Metric: 0.80 |
INFO:pytorch_lightning.callbacks.early_stopping:Metric val_metric improved by
0.236 >= min_delta = 0.0. New best score: 0.800
Train_Loss: 0.44, Train_Metric: 0.85
Validation: 0it [00:00, ?it/s]
Epoch 3: Val_Loss: 0.79, Val_Metric: 0.71 | Train_Loss: 0.37, Train_Metric: 0.87
Epoch 00003: reducing learning rate of group 0 to 1.0000e-03.
Validation: 0it [00:00, ?it/s]
Epoch 4: Val_Loss: 0.50, Val_Metric: 0.86 |
INFO:pytorch_lightning.callbacks.early_stopping:Metric val_metric improved by
0.055 >= min_delta = 0.0. New best score: 0.855
Train_Loss: 0.30, Train_Metric: 0.90
Validation: 0it [00:00, ?it/s]
Epoch 5: Val_Loss: 1.67, Val_Metric: 0.47 | Train_Loss: 0.29, Train_Metric: 0.91
Epoch 00005: reducing learning rate of group 0 to 5.0000e-04.
Validation: 0it [00:00, ?it/s]
Epoch 6: Val_Loss: 0.71, Val_Metric: 0.77 | Train_Loss: 0.25, Train_Metric: 0.92
Epoch 00006: reducing learning rate of group 0 to 2.5000e-04.
Validation: 0it [00:00, ?it/s]
Epoch 7: Val_Loss: 0.30, Val_Metric: 0.91 |
INFO:pytorch_lightning.callbacks.early_stopping:Metric val_metric improved by
0.055 >= min_delta = 0.0. New best score: 0.910
Train_Loss: 0.21, Train_Metric: 0.93
Validation: 0it [00:00, ?it/s]
Epoch 8: Val_Loss: 0.55, Val_Metric: 0.82 | Train_Loss: 0.20, Train_Metric: 0.94
Epoch 00008: reducing learning rate of group 0 to 1.2500e-04.
Validation: 0it [00:00, ?it/s]
Epoch 9: Val_Loss: 0.30, Val_Metric: 0.90 | Train_Loss: 0.17, Train_Metric: 0.95
Epoch 00009: reducing learning rate of group 0 to 6.2500e-05.
Validation: 0it [00:00, ?it/s]
Epoch 10: Val_Loss: 0.21, Val_Metric: 0.93 |

```

```
INFO:pytorch_lightning.callbacks.early_stopping:Metric val_metric improved by
0.019 >= min_delta = 0.0. New best score: 0.928
INFO:pytorch_lightning.utilities.rank_zero:`Trainer.fit` stopped:
`max_epochs=10` reached.

Train_Loss: 0.15, Train_Metric: 0.96
```

```
[22]: file = f"{trainer.logger.log_dir}/metrics.csv"
      plot_losses_acc(file)
```



```
[23]: ckpt_path = trainer.checkpoint_callback.best_model_path
      train_acc = trainer.validate(dataloaders=dm.train_dataloader(),
      ↪ckpt_path=ckpt_path, verbose=False)
      valid_acc = trainer.validate(dataloaders=dm.val_dataloader(),
      ↪ckpt_path=ckpt_path, verbose=False)
      print(f"Train Accuracy: {train_acc[0]['val_metric']*100:0.2f}")
      print(f"Validation Accuracy: {valid_acc[0]['val_metric']*100:0.2f}")
      wandb.finish()
```

```
INFO:pytorch_lightning.utilities.rank_zero:Restoring states from the checkpoint
path at /content/drive/MyDrive/data/models/dl_fall_2023/fmnist/oct-
31/logs/epoch=9-step=3750-v3.ckpt
INFO:pytorch_lightning.accelerators.cuda:LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES:
[0]
INFO:pytorch_lightning.utilities.rank_zero:Loaded model weights from the
checkpoint at /content/drive/MyDrive/data/models/dl_fall_2023/fmnist/oct-
31/logs/epoch=9-step=3750-v3.ckpt
/usr/local/lib/python3.10/dist-
packages/pytorch_lightning/trainer/connectors/data_connector.py:490:
PossibleUserWarning: Your `val_dataloader`'s sampler has shuffling enabled, it
```



```

is strongly recommended that you turn shuffling off for val/test dataloaders.
    rank_zero_warn(
/usr/local/lib/python3.10/dist-
packages/pytorch_lightning/trainer/connectors/data_connector.py:442:
PossibleUserWarning: The dataloader, val_dataloader, does not have many workers
which may be a bottleneck. Consider increasing the value of the `num_workers`
argument` (try 8 which is the number of cpus on this machine) in the
`DataLoader` init to improve performance.
    rank_zero_warn(

Validation: 0it [00:00, ?it/s]

Epoch 11: Val_Loss: 0.14, Val_Metric: 0.96 |

INFO:pytorch_lightning.utilities.rank_zero:Restoring states from the checkpoint
path at /content/drive/MyDrive/data/models/dl_fall_2023/fmnist/oct-
31/logs/epoch=9-step=3750-v3.ckpt
INFO:pytorch_lightning.accelerators.cuda:LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES:
[0]
INFO:pytorch_lightning.utilities.rank_zero:Loaded model weights from the
checkpoint at /content/drive/MyDrive/data/models/dl_fall_2023/fmnist/oct-
31/logs/epoch=9-step=3750-v3.ckpt

Validation: 0it [00:00, ?it/s]

Epoch 11: Val_Loss: 0.21, Val_Metric: 0.93 | Train Accuracy: 96.32
Validation Accuracy: 92.81

<IPython.core.display.HTML object>

VBox(children=(Label(value='0.006 MB of 0.006 MB uploaded (0.000 MB_
deduped)\r'), FloatProgress(value=1.0, max...

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

```